Annotated Literature Companion

for "Babel's Machines: Linguistic Relativity, Cooperation, and the Emerging Behaviour of Multilingual Artificial Agents"

Compiled: 2026-04-27

Foreword

This companion document explains every paper cited in *Babel's Machines*. For each entry we provide: the bib key, full citation, a **verifiable source URL**, a one-paragraph description of what the paper does, and one paragraph on how the present perspective uses, extends, contradicts, or builds on it.

Verification. All cited entries were independently verified against arXiv, ACL Anthology, journal websites (Nature, Science, PNAS, Royal Society, Cambridge Core, IOS Press, MIT Press), and OpenReview. The full audit trail is in `docs/reviews/citation_audit_round1.md`. Reviewer Round 1 and Round 2 additions were web-verified at insertion time. Three categories were distinguished:

- **Verified clean (60+ entries):** title, authors, year, venue, and arXiv ID match authoritative sources.
- **Fixed during audit (12 entries):** discrepancies (wrong venue, missing authors, wrong title, wrong year) were corrected. The most consequential were `han2022pledges` (venue was incorrectly listed as *Royal Society Open Science*; corrected to ALIFE 2022) and `turner2023actadd` (wrong title and missing author).
- **Cannot fully verify (1 entry):** `repello2024multilingualguardrails` is a blog post; the title and content match https://repello.ai/blog/, but date metadata could not be independently confirmed.

Entries are grouped thematically, not alphabetically, so the reader can see the structure of the literature the paper draws on.

Contents

1 Foundations: classical linguistic relativity and the cognitive baseline **1**

2 The machine-behaviour programme **2**

3 Cooperation: the evolutionary game-theory canon **3**

4 Cultural evolution and the WEIRD problem **4**

5 Theory of mind: foundations and modern critique **5**

6 Foundational LLM architecture and training **6**

7 Multilingual representation: classics 7

8 Mechanistic interpretability: hub, barrier, and alternatives 7

9 Interpretability: causal-evidence precedents (Round 2 additions) 8

10 Calibration across languages 9

11 Multilingual safety, jailbreaks, and adversarial probes 10

12 Multilingual social-dilemma evidence 11

13 Moderating factors beyond language (Round 1 additions) 13

14 Negotiation, persuasion, and auction 14

15 Social-deduction games 15

16 Cooperative communication and theory-of-mind games 15

17 Open-world game environments 16

18 Dialogue games (clembench) 17

19 Multi-agent benchmarks and failure analysis 17

20 Pretraining-corpus audits 18

21 Other supporting citations 18

1 Foundations: classical linguistic relativity and the cognitive baseline

`boroditsky2001does` — Boroditsky, L. (2001). Does language shape thought? Mandarin and English speakers' conceptions of time. *Cognitive Psychology* 43(1):1–22.
Verify: https://doi.org/10.1006/cogp.2001.0748
What it is. Classical experimental study showing that Mandarin and English speakers conceptualise time on different spatial axes (vertical vs. horizontal), depending on the metaphors built into their respective languages.
How Babel's Machines uses it. We cite this as the modern, measured form of the linguistic-relativity hypothesis. The Algorithmic Sapir–Whorf Switch is named after this tradition; we explicitly distance ourselves from the strong determinist Whorfian reading and adopt Boroditsky's "cognitive affordances" framing.

`boroditsky2011languageshapes` — Boroditsky, L. (2011). How language shapes thought. *Scientific American* 304(2):62–65.
Verify: https://doi.org/10.1038/scientificamerican0211-62
What it is. Popular-science synthesis of two decades of evidence for measured linguistic relativity, covering colour categorisation, spatial orientation in Kuuk Thaayorre, and grammatical gender effects.
How Babel's Machines uses it. Used as the second pillar of our linguistic-relativity background. Together with Boroditsky 2001, it grounds our claim that "language as cognitive affordance" is a credible framework to translate to machine populations.

2 The machine-behaviour programme

`rahwan2019machinebehaviour` — Rahwan, I., Cebrian, M., Obradovich, N., et al. (2019). Machine behaviour. *Nature* 568(7753):477–486.
Verify: https://doi.org/10.1038/s41586-019-1138-y
What it is. Manifesto-style *Nature* paper proposing that AI systems be studied as a class of behaving agents using tools from biology, economics, and social science — not solely the engineering discipline that built them.
How Babel's Machines uses it. This is the anchor of our entire argument. The opening of the paper takes Rahwan's programme as the point of departure; we frame the perspective as a contribution within it. Iyad Rahwan is also a guest editor of the theme issue.

`han2022emergent` — Han, T.A. (2022). Emergent behaviours in multi-agent systems with evolutionary game theory. *AI Communications* 35(4):327–337.
Verify: https://doi.org/10.3233/AIC-220104
What it is. Argues that evolutionary game theory provides the natural lens for studying emergent behaviour in multi-agent systems, including those whose strategies arise from learning rather than mutation.
How Babel's Machines uses it. We cite this as the methodological bridge between machine populations and evolutionary game theory. Han is a guest editor of the theme issue; the citation is intentional and substantive.

`han2022pledges` — Han, T.A., Santos, F.C., Pereira, L.M., Lenaerts, T. (2022). Voluntary safety pledges overcome over-regulation dilemma in AI development. *Proc. ALIFE 2022.*
Verify: https://doi.org/10.1162/isal_a_00484
What it is. Models AI-development regulation as an evolutionary game; shows voluntary pledges with peer or institutional sanctions can resolve the over-regulation dilemma.
How Babel's Machines uses it. Cited in our governance discussion as the formal precedent for treating multilingual-AI deployment regulation as an evolutionary-game-theoretic problem. **Note:** the original bib entry incorrectly listed *Royal Society Open Science* as the venue; corrected to ALIFE 2022 during the audit.

3 Cooperation: the evolutionary game-theory canon

`axelrod1984evolution` — Axelrod, R. (1984). *The Evolution of Cooperation.* Basic Books, New York.
Verify: https://www.basicbooks.com/titles/robert-axelrod/the-evolution-of-cooperation/9780465005642/
What it is. The foundational book demonstrating, via computer tournaments, that cooperative strategies (Tit-for-Tat) can emerge and stabilise among self-interested agents through repeated interaction.
How Babel's Machines uses it. We use Axelrod as the negative anchor for the "culturally inherited cooperation" argument: the Axelrodian view says cooperation is achieved through repeated interaction and reputation; LLMs cooperate *before* any of those scaffolds operate.

`nowak2006five` — Nowak, M.A. (2006). Five rules for the evolution of cooperation. *Science* 314(5805):1560–1563.
Verify: https://doi.org/10.1126/science.1133755
What it is. Synthesises the five mechanisms by which cooperation can evolve: kin selection, direct reciprocity, indirect reciprocity, network reciprocity, group selection.
How Babel's Machines uses it. Cited as the canonical statement of the mechanisms that LLM

cooperation does *not* require, supporting our claim that LLM prosocial defaults are a different kind of phenomenon entirely — cultural inheritance, not evolved reciprocity.

`nowak1998evolution` — Nowak, M.A., Sigmund, K. (1998). Evolution of indirect reciprocity by image scoring. *Nature* 393(6685):573–577.
Verify: https://doi.org/10.1038/31225
What it is. Establishes that reputation-based "image scoring" enables cooperation to evolve when direct reciprocity fails.
How Babel's Machines uses it. One of the canonical scaffolds whose absence is consequential for the LLM case: LLMs cooperate without reputation, image scoring, or repeated interaction. The contrast sharpens the cultural-inheritance reading.

`panchanathan2004indirect` — Panchanathan, K., Boyd, R. (2004). Indirect reciprocity can stabilize cooperation without the second-order free rider problem. *Nature* 432(7016):499–502.
Verify: https://doi.org/10.1038/nature02978
What it is. Shows that indirect reciprocity, properly modelled, escapes the second-order free-rider problem and stabilises cooperation in larger populations.
How Babel's Machines uses it. Cited in the evolutionary section as another scaffold whose absence in LLM populations the cultural-inheritance reading must explain.

`perc2017statistical` — Perc, M., Jordan, J.J., Rand, D.G., Wang, Z., Boccaletti, S., Szolnoki, A. (2017). Statistical physics of human cooperation. *Physics Reports* 687:1–51.
Verify: https://doi.org/10.1016/j.physrep.2017.05.004
What it is. Comprehensive review of the statistical-physics approach to cooperation: spatial reciprocity, network reciprocity, phase transitions, and the dynamics of cooperative regimes.
How Babel's Machines uses it. Cited as the second pillar of the theoretical machinery we propose to apply to LLM populations. Matǰaz Perc is a guest editor of the theme issue; the link to his programme is intentional.

`neumann1944theory` — von Neumann, J., Morgenstern, O. (1944). *Theory of Games and Economic Behavior*. Princeton University Press.
Verify: https://press.princeton.edu/books/paperback/9780691130613/theory-of-games-and-economic-behavior
What it is. The book that founded modern game theory as a discipline.
How Babel's Machines uses it. Foundational acknowledgement only; not centrally argued from.

`nash1950equilibrium` — Nash, J.F. (1950). Equilibrium points in n-person games. *Proc. National Academy of Sciences* 36(1):48–49.
Verify: https://doi.org/10.1073/pnas.36.1.48
What it is. The two-page note that introduced the Nash equilibrium concept.
How Babel's Machines uses it. Foundational acknowledgement only.

4 Cultural evolution and the WEIRD problem

`boyd1985culture` — Boyd, R., Richerson, P.J. (1985). *Culture and the Evolutionary Process*. University of Chicago Press.
Verify: https://press.uchicago.edu/ucp/books/book/chicago/C/bo3621168.html
What it is. Foundational work formalising cultural evolution as a Darwinian process operating on transmitted information, with mechanisms (oblique/horizontal/vertical inheritance, conformist transmission, etc.) analogous to genetic evolution.
How Babel's Machines uses it. Central to our reframe. The claim that pretraining + RLHF act

as cultural inheritance is explicitly framed in Boyd–Richerson terms; LLMs are cast as a new substrate for their framework.

`richerson2005notbygenes` — Richerson, P.J., Boyd, R. (2005). *Not by Genes Alone: How Culture Transformed Human Evolution*. University of Chicago Press.
Verify: https://press.uchicago.edu/ucp/books/book/chicago/N/bo3636224.html
What it is. Companion volume to Boyd & Richerson 1985, accessible to a wider audience, summarising the case for cumulative cultural evolution.
How Babel's Machines uses it. Cited together with Boyd 1985 to anchor the cultural-inheritance reading.

`henrich2016secret` — Henrich, J. (2015). *The Secret of Our Success: How Culture Is Driving Human Evolution....* Princeton University Press.
Verify: https://press.princeton.edu/books/hardcover/9780691166858/the-secret-of-our-success
What it is. Argues that humans are an ultra-cooperative species because of cumulative culture, not genes; cooperation is built up by transmission, not selected for biologically.
How Babel's Machines uses it. Used as the third anchor of the cultural-inheritance argument. Henrich's transmission framing is what we apply to machine-text-trained agents.

`henrich2010weirdest` — Henrich, J., Heine, S.J., Norenzayan, A. (2010). The weirdest people in the world? *Behavioral and Brain Sciences* 33(2–3):61–83.
Verify: https://doi.org/10.1017/S0140525X0999152X
What it is. Shows that Western, Educated, Industrialised, Rich, Democratic (WEIRD) populations are statistical outliers on many traits the cooperation literature treats as universal.
How Babel's Machines uses it. This is the linchpin of the WEIRD-corollary argument. If LLM cooperation is culturally inherited from WEIRD-skewed corpora, LLMs encode WEIRD norms while presenting as universal — a normative externality of multilingual deployment.

`muthukrishna2020beyondweird` — Muthukrishna, M., Bell, A.V., Henrich, J., Curtin, C.M., Gedranovich, A., McInerney, J., Thue, B. (2020). Beyond WEIRD Psychology. *Psychological Science* 31(6):678–701.
Verify: https://doi.org/10.1177/0956797620916782
What it is. Operationalises "WEIRD" as a measurable cultural distance and maps where in the world WEIRDness deviates from local norms.
How Babel's Machines uses it. Cited together with Henrich 2010 to make the WEIRD argument quantitative rather than rhetorical. Provides the methodology to test whether LLM defaults track WEIRD distance.

`tomasello2009whywecooperate` — Tomasello, M. (2009). *Why We Cooperate*. MIT Press.
Verify: https://mitpress.mit.edu/9780262013598/why-we-cooperate/
What it is. Developmental and comparative argument that human cooperation rests on shared intentionality and joint attention — cognitive primitives largely absent in other apes.
How Babel's Machines uses it. Cited as part of the broader cooperation-and-cognition canon that the perspective draws on; supports the claim that ToM is a cognitive prerequisite for cooperation.

5 Theory of mind: foundations and modern critique

`premack1978tom` — Premack, D., Woodruff, G. (1978). Does the chimpanzee have a theory of mind? *Behavioral and Brain Sciences* 1(4):515–526.
Verify: https://doi.org/10.1017/S0140525X00076512

What it is. The 1978 paper that named "theory of mind" as a research target in comparative cognition.
How Babel's Machines uses it. Cited as the etymological anchor when we introduce ToM in Section 5.

`deweerd2013bayesiantom` — de Weerd, H., Verbrugge, R., Verheij, B. (2013). How much does it help to know what she knows you know? *Artificial Intelligence* 199–200:67–92.
Verify: https://doi.org/10.1016/j.artint.2013.05.004
What it is. Agent-based simulation establishing that recursive ToM (k-level reasoning) provides diminishing returns past first or second order, with first-order being the strongest predictor of cooperative success.
How Babel's Machines uses it. Cited as the formal ground for the LLM-Hanabi finding that *first-order* ToM is what predicts cooperation in LLM agents — consistent with the human/AI agent literature.

`byrne1988machiavellian` — Byrne, R.W., Whiten, A. (eds.) (1988). *Machiavellian Intelligence.* Oxford University Press.
Verify: https://global.oup.com/academic/product/machiavellian-intelligence-9780198521754
What it is. The volume that established the "Machiavellian intelligence" hypothesis: that primate cognitive evolution was driven by the demands of social manipulation, not domain-general problem-solving.
How Babel's Machines uses it. Cited in the SPIN-Bench discussion: the planning-vs-social-intelligence dissociation in LLMs mirrors the evolutionary distinction between domain-general planning and Machiavellian social cognition.

`ullman2023tomfail` — Ullman, T. (2023). Large Language Models Fail on Trivial Alterations to Theory-of-Mind Tasks. arXiv:2302.08399.
Verify: https://arxiv.org/abs/2302.08399
What it is. Demonstrates that small variations to standard false-belief items collapse apparent LLM ToM performance, suggesting reported ToM successes may be surface pattern-matching rather than genuine mental-state representation.
How Babel's Machines uses it. Round 2 reviewer addition. Cited explicitly in the ToM construct-validity caveat: even if absolute LLM ToM scores are contestable, our claim about *cross-lingual differences* on the same items remains well-defined.

`sap2022neuraltom` — Sap, M., LeBras, R., Fried, D., Choi, Y. (2022). Neural Theory-of-Mind? On the Limits of Social Intelligence in Large LMs. *EMNLP 2022.* arXiv:2210.13312.
Verify: https://arxiv.org/abs/2210.13312
What it is. Reports that GPT-3, ChatGPT, and GPT-4 reach only $\sim 60\%$ on SocialIQa and ToMi benchmarks, well below human; argues scaling alone is insufficient for ToM.
How Babel's Machines uses it. Round 2 reviewer addition. Pairs with Ullman 2023 in the construct-validity caveat. The combined citation acknowledges the controversy without conceding the cross-lingual asymmetry argument.

`kovacs2009bilingual` — Kovács, Á.M. (2009). Early bilingualism enhances mechanisms of false-belief reasoning. *Developmental Science* 12(1):48–54.
Verify: https://doi.org/10.1111/j.1467-7687.2008.00742.x
What it is. Shows that bilingual children pass false-belief tasks earlier than monolingual peers, linking the executive demands of code-switching to ToM.
How Babel's Machines uses it. This is the human counter-prediction that frames the LLM ToM

result as surprising: in humans, multilingualism *strengthens* ToM; in LLMs, it weakens it. The contrast is rhetorical and substantive.

`goetz2003bilingual` — Goetz, P.J. (2003). The effects of bilingualism on theory of mind development. *Bilingualism: Language and Cognition* 6(1):1–15.
Verify: https://doi.org/10.1017/S1366728903001007
What it is. Independent confirmation of the bilingual ToM advantage with a different developmental population.
How Babel's Machines uses it. Paired with Kovács 2009 to make the bilingual-ToM finding robust rather than a one-study artefact.

6 Foundational LLM architecture and training

`vaswani2017transformer` — Vaswani, A., Shazeer, N., Parmar, N., et al. (2017). Attention is all you need. *NeurIPS 2017.*
Verify: https://arxiv.org/abs/1706.03762
What it is. The original Transformer architecture paper; introduced self-attention as the dominant primitive for sequence modelling.
How Babel's Machines uses it. Cited as the architectural substrate; everything we discuss assumes Transformer-based LLMs.

`brown2020gpt3` — Brown, T., Mann, B., Ryder, N., et al. (2020). Language models are few-shot learners. *NeurIPS 2020.*
Verify: https://arxiv.org/abs/2005.14165
What it is. The GPT-3 paper; showed few-shot in-context learning emerges at scale and that English-dominant pretraining produces an asymmetric multilingual model.
How Babel's Machines uses it. Cited when we note that pretraining corpora are largely English-dominant; this is the statistical asymmetry that the Semantic Hub Hypothesis later explains.

`ouyang2022training` — Ouyang, L., Wu, J., Jiang, X., et al. (2022). Training language models to follow instructions with human feedback. *NeurIPS 2022.*
Verify: https://arxiv.org/abs/2203.02155
What it is. The InstructGPT paper; established RLHF as a standard alignment recipe.
How Babel's Machines uses it. Cited where we discuss RLHF as the second leg of cultural inheritance and as a documented source of linguistic skew (RLHF annotators are disproportionately English-speaking).

`bai2022constitutional` — Bai, Y., Kadavath, S., Kundu, S., et al. (2022). Constitutional AI: Harmlessness from AI feedback. arXiv:2212.08073.
Verify: https://arxiv.org/abs/2212.08073
What it is. Anthropic's RLAIF approach: replaces human preference labels with AI-generated critiques against a written "constitution".
How Babel's Machines uses it. Cited alongside Ouyang to broaden the RLHF discussion to include AI-feedback variants; the linguistic-skew problem persists in RLAIF if the critic model is itself English-skewed.

7 Multilingual representation: classics

`conneau2020xlmr` — Conneau, A., Khandelwal, K., Goyal, N., et al. (2020). Unsupervised cross-lingual representation learning at scale (XLM-R). *ACL 2020.*

Verify: https://arxiv.org/abs/1911.02116
What it is. Trained a multilingual masked language model on 100 languages and 2.5TB of CommonCrawl; established that scale + balanced sampling can rival monolingual models on individual languages.
How Babel's Machines uses it. Round 1 reviewer addition. Cited in the architectural-alternatives subsection (§4.1): scale alleviates but does not eliminate cross-lingual subspace interference.

`xue2021mt5` — Xue, L., Constant, N., Roberts, A., et al. (2021). mT5: A massively multilingual pre-trained text-to-text transformer. *NAACL 2021.*
Verify: https://arxiv.org/abs/2010.11934
What it is. Multilingual T5 across 101 languages; documented "accidental translation" (the model producing output in unintended languages) as a failure mode.
How Babel's Machines uses it. Round 1 reviewer addition. Cited together with XLM-R as evidence that subspace interference persists even in well-engineered multilingual models. Also cited as an example of a calibration-repair lever (language adapters / instruction-tuning variations).

`pires2019multilingualbert` — Pires, T., Schlinger, E., Garrette, D. (2019). How multilingual is multilingual BERT? *ACL 2019.*
Verify: https://arxiv.org/abs/1906.01502
What it is. Showed that multilingual BERT supports surprising zero-shot cross-lingual transfer but exhibits "systematic deficiencies" on certain language pairs.
How Babel's Machines uses it. Round 1 reviewer addition. The "systematic deficiencies" finding is the empirical base for the language-specific microcircuits alternative to the Semantic Hub.

8 Mechanistic interpretability: hub, barrier, and alternatives

`wu2024semantichub` — Wu, Z., Yu, X.V., Yogatama, D., Lu, J., Kim, Y. (2024/2025). The Semantic Hub Hypothesis. *ICLR 2025.*
Verify: https://arxiv.org/abs/2411.04986
What it is. Provides causal evidence (via activation steering) that multilingual LLMs project heterogeneous linguistic inputs into a centralised, English-dominant intermediate hub.
How Babel's Machines uses it. This is one of the two architectural pillars of Section 4. We cite it for both the representational claim and the activation-steering methodology that we propose to extend with a behavioural readout in P2.

`wendler2024llamasenglish` — Wendler, C., Veselovsky, V., Monea, G., West, R. (2024). Do Llamas Work in English? On the Latent Language of Multilingual Transformers. *ACL 2024.*
Verify: https://aclanthology.org/2024.acl-long.820/
What it is. Logit-lens evidence that Llama-family decoders pass through an intermediate state that is statistically closer to English than to the input language.
How Babel's Machines uses it. Cited as the "latent-language priors" alternative to the Semantic Hub — a stronger claim that the hub is specifically English-anchored, not just shared. The natural target for the depth-resolved probes proposed in P2.

`bafna2025translationbarrier` — Bafna, N., Li, T., Murray, K., Mortensen, D.R., Yarowsky, D., Sirin, H., Khashabi, D. (2025). The Translation Barrier Hypothesis. *IJCNLP-AACL 2025.*
Verify: https://arxiv.org/abs/2506.22724
What it is. Shows that many multilingual generation failures are not upstream-reasoning failures but late-layer realisation failures: the model reaches the right conclusion internally and then fails projecting into the target language.

How Babel's Machines uses it. The second architectural pillar. Combined with the Semantic Hub, gives the hub-vs-realisation diagnostic categories that P2 operationalises.

`turner2023actadd` — Turner, A.M., Thiergart, L., Leech, G., Udell, D., Vazquez, J.J., Mini, U., MacDiarmid, M. (2023). Steering Language Models With Activation Engineering. arXiv:2308.10248.
Verify: https://arxiv.org/abs/2308.10248
What it is. Introduced Activation Addition (ActAdd) as a lightweight, training-free way to steer LLM behaviour by adding activation differences at inference time.
How Babel's Machines uses it. Methodological precedent for the steering-vector experiment in P2. **Note:** the original bib entry had the wrong title ("Activation Addition: Steering Language Models Without Optimization") and a missing author (Vazquez); both fixed during the audit.

`kumar2025indictunedlens` — Panchal, M., Varshney, D., Mamta, Ekbal, A. (2025/2026). Indic-TunedLens: Interpreting Multilingual Models in Indian Languages. *EACL 2026 VarDial Workshop*.
Verify: https://arxiv.org/abs/2602.15038
What it is. Adapts the TunedLens technique to Indic languages, providing per-language layer-resolved interpretability tooling for low-resource scripts.
How Babel's Machines uses it. Cited as the methodology for distinguishing hub failures from realisation failures in P2; without per-language probes, the hub/barrier diagnostic cannot be operationalised in low-resource languages.

9 Interpretability: causal-evidence precedents (Round 2 additions)

`chen2023spatialcausal` — Chen, Y., Gan, Y., Li, S., Yao, L., Zhao, X. (2023). More than Correlation: Do Large Language Models Learn Causal Representations of Space? arXiv:2312.16257.
Verify: https://arxiv.org/abs/2312.16257
What it is. Causal-intervention experiments showing that spatial representations in DeBERTa and GPT-Neo *causally* affect next-token prediction, not merely correlate with it.
How Babel's Machines uses it. Round 2 addition. Cited in §4.3 (Closing the Causal Loop) as the methodological precedent for the steering-vector → behaviour experiment in P2.

`zhang2026worldmodels` — Zhang, X. (2026). What Do World Models Learn in RL? Probing Latent Representations in Learned Environment Simulators. arXiv:2603.21546.
Verify: https://arxiv.org/abs/2603.21546
What it is. Demonstrates that latent state variables in IRIS and DIAMOND world models are linearly decodable *and* causally engaged via targeted intervention.
How Babel's Machines uses it. Round 2 addition. Used as a second causal-intervention precedent that justifies the move from representational diagnosis to behavioural readout in P2.

`yan2025arithmetic` — Yan, Y. (2025). Addition in Four Movements: Mapping Layer-wise Information Trajectories in LLMs. arXiv:2506.07824.
Verify: https://arxiv.org/abs/2506.07824
What it is. Identifies a coherent four-stage layer-by-layer trajectory in Llama-3 doing multi-digit addition: formula representation → numerical features → digit-by-digit results → final token.
How Babel's Machines uses it. Round 2 addition. Cited as evidence that depth-resolved probing can recover meaningful processing stages, justifying the L/3 and 2L/3 probe choices in P2.

`lee2024causalpaths` — Lee, I., Lum, J., Liu, Z., Yogatama, D. (2024). Causal Interventions on Causal Paths: Mapping GPT-2's Reasoning From Syntax to Semantics. arXiv:2410.21353.

Verify: https://arxiv.org/abs/2410.21353
What it is. Causal intervention study showing GPT-2 processes causal-language sentences via early-layer syntactic detection followed by late-layer semantic analysis.
How Babel's Machines uses it. Round 2 addition. Cited as direct precedent for the early-layers-vs-late-layers dissociation that P2's hub-vs-realisation split assumes.

`zhenyu2025logitlens` — Hong, Z. (2025). LogitLens4LLMs: Extending Logit Lens Analysis to Modern Large Language Models. arXiv:2503.11667.
Verify: https://arxiv.org/abs/2503.11667
What it is. Open-source toolkit extending the Logit Lens technique to Llama-3.1, Qwen-2.5, and other modern architectures, with memory-efficient hooks for layer-wise analysis up to 70B parameters.
How Babel's Machines uses it. Round 2 addition. Cited as the practical tooling that makes P2 immediately runnable on frontier models.

10 Calibration across languages

`ahuja2022calibration` — Ahuja, K., Sitaram, S., Dandapat, S., Choudhury, M. (2022). On the Calibration of Massively Multilingual Language Models. *EMNLP 2022.*
Verify: https://aclanthology.org/2022.emnlp-main.290/
What it is. Documents the cross-lingual calibration gap on the MEGA benchmark across XGLM, BLOOM, and PaLM with 15-bin equal-width binning; reports ECE rising from $\sim 4.6\%$ in English to $\sim 23.1\%$ in non-English.
How Babel's Machines uses it. The empirical source for the "$\sim 5\times$ ECE gap" claim that anchors the calibration-crisis discussion in Section 4.

`zhou2025beyond` — Zhou, E., Zhang, C., Hu, T., Li, C., Collier, N., Vulić, I., Korhonen, A. (2025). Beyond the Final Layer: Intermediate Representations for Better Multilingual Calibration. arXiv:2510.03136.
Verify: https://arxiv.org/abs/2510.03136
What it is. Reproduces the cross-lingual ECE gap on Llama-3 and Qwen-2 and shows that probing a late-intermediate layer (rather than the final layer) recovers part of the gap.
How Babel's Machines uses it. Cited together with Ahuja 2022 to show the ECE asymmetry replicates across model generations and to motivate intermediate-layer probing in P2.

`fu2025multilingualjudge` — Fu, X., Liu, W. (2025). How Reliable is Multilingual LLM-as-a-Judge? *Findings of EMNLP 2025.*
Verify: https://aclanthology.org/2025.findings-emnlp.587/
What it is. Empirically establishes that single-judge LLM-as-Judge reliability collapses to Fleiss' $\kappa \approx 0.3$ across languages.
How Babel's Machines uses it. Direct empirical anchor for P4 (Evaluation reliability ceiling). The whole prediction is structured around the question of whether ensemble protocols can lift this ceiling above 0.5.

11 Multilingual safety, jailbreaks, and adversarial probes

`yong2024lowresource` — Yong, Z.-X., Menghini, C., Bach, S.H. (2023). Low-Resource Languages Jailbreak GPT-4. *NeurIPS 2023 SoLaR Workshop* (Best Paper).
Verify: https://arxiv.org/abs/2310.02446

What it is. Demonstrates that translating prompts into low-resource languages (Zulu, Scots Gaelic, Hmong, Guarani) bypasses GPT-4's safety training.
How Babel's Machines uses it. Cited as evidence for the "RLHF safety coverage is skewed" confound in Section 3.3, and for the broader claim that linguistic asymmetries open attack surfaces.

`yong2025multilingualsafety` — Yong, Z.-X., Ermis, B., Fadaee, M., Bach, S.H., Kreutzer, J. (2025). The State of Multilingual LLM Safety Research: From Measuring the Language Gap to Mitigating It. *EMNLP 2025.*
Verify: https://aclanthology.org/2025.emnlp-main.800/
What it is. Survey of the multilingual safety literature; documents the gap between English and non-English safety coverage and proposes mitigation directions.
How Babel's Machines uses it. Cited as the modern survey anchor for the "RLHF coverage is skewed" claim and as one of the safety benchmarks to be paired with interactive multi-agent ones in our governance recommendations.

`yoo2025csrt` — Yoo, H., Yang, Y., Lee, H. (2025). Code-Switching Red-Teaming. *ACL 2025.*
Verify: https://aclanthology.org/2025.acl-long.657/
What it is. Shows that intra-utterance language mixing achieves a 46.7% increase in attack success against frontier LLM safety systems.
How Babel's Machines uses it. The empirical source for the code-switching attack-surface claim. Cited in Section 7 for societal consequences and again in the governance levers (Code-Switching probes as default audit).

`he2025agentmiddle` — He, P., Lin, Y., Dong, S., Xu, H., Xing, Y., Liu, H. (2025). Red-Teaming LLM Multi-Agent Systems via Communication Attacks. *Findings of ACL 2025.*
Verify: https://aclanthology.org/2025.findings-acl.349/
What it is. Introduces Agent-in-the-Middle (AiTM) attacks that intercept and manipulate inter-agent messages in multilingual multi-agent systems.
How Babel's Machines uses it. Cited together with CSRT as the multi-agent extension of the attack-surface argument; both become default audit probes in our governance recommendations.

`m-alert2025` — Friedrich, F., Tedeschi, S., Schramowski, P., Brack, M., Navigli, R., Nguyen, H., Li, B., Kersting, K. (2025). LLMs Lost in Translation: M-ALERT uncovers Cross-Linguistic Safety Inconsistencies. *ICLR 2025.*
Verify: https://arxiv.org/abs/2412.15035
What it is. Multilingual safety benchmark documenting cross-linguistic inconsistency in alignment behaviour.
How Babel's Machines uses it. Cited as one of the safety benchmarks to be paired with interactive multi-agent benchmarks in our diversification recommendation. **Note:** the original bib entry had wrong title and venue ("OpenReview"); both corrected during the audit.

`wang2024evaluating` — Huang, K., Mo, F., Zhang, X., Li, H., Li, Y., Zhang, Y., Yi, W., Mao, Y., Liu, J., Xu, Y., Xu, J., Nie, J.-Y., Liu, Y. (2024). A Survey on Large Language Models with Multilingualism: Recent Advances and New Frontiers. arXiv:2405.10936.
Verify: https://arxiv.org/abs/2405.10936
What it is. Comprehensive survey of multilingual LLM research, covering pretraining, alignment, evaluation, and deployment gaps.
How Babel's Machines uses it. Cited as broader survey context in Section 7. **Note:** earlier bib entry was missing author Xinyu Zhang; corrected during the audit.

`repello2024multilingualguardrails` — Repello AI (2024). Introducing New Multilingual AI Safety Guardrails for 100 Languages. Blog post.

Verify: https://repello.ai/blog/introducing-new-multilingual-ai-safety-guardrails-for-100-languages
What it is. Industry blog post announcing a multilingual safety guardrail system covering 100 languages.
How Babel's Machines uses it. Cited as evidence that the industry is starting to address the multilingual safety gap. **Note:** this is the one citation the audit could not fully verify (date metadata not confirmable); the user may want to substitute with the academic CREST paper.

`stanfordhai2024languagegap` — Stanford HAI Policy Brief (2025). Mind the (Language) Gap: Mapping the Challenges of LLM Development in Low-Resource Language Contexts.
Verify: https://hai.stanford.edu/policy/mind-the-language-gap
What it is. Stanford HAI policy brief mapping the structural challenges of LLM development for low-resource languages.
How Babel's Machines uses it. Cited in Section 7 as policy-community recognition of the language gap as a structural issue. **Note:** year corrected from 2024 to 2025 during the audit.

12 Multilingual social-dilemma evidence

`huynh2026moreatstake` — Huynh, T.-K., Dao-Sy, D.-M., Cao, T.-B., et al. (2026). More at Stake: How Payoff and Language Shape LLM Agent Strategies in Cooperation Dilemmas. arXiv:2601.19082.
Verify: https://arxiv.org/abs/2601.19082
What it is. Large-scale multilingual IPD evaluation across four model families and five languages, with strategy-classifier labelling at varying confidence thresholds.
How Babel's Machines uses it. The primary empirical source for Table 1. Most of the cross-lingual divergence patterns we discuss in Section 3.1 come from this study (the user's own work).

`huynh2025fairgame` — Huynh, T.-K., Dao-Sy, D.-M., Cao, T.-B., et al. (2025). Understanding LLM Agent Behaviours via Game Theory. arXiv:2512.07462.
Verify: https://arxiv.org/abs/2512.07462
What it is. Companion methodological paper to the IPD study; details strategy-recognition methodology, multi-agent dynamics, and bias analysis.
How Babel's Machines uses it. Cited together with the More-at-Stake paper for methodological details, and specifically for the Vietnamese low-stakes "defection" = arithmetic-error finding.

`za2026towards` — Za, J., Panos, A., Cuhel, J., Albanie, S. (2026). Towards Predictive Models of Strategic Behaviour in Large Language Model Agents. *ICLR 2026 Workshop*.
Verify: https://openreview.net/forum?id=cBzy30XXm5
What it is. Builds predictive models of LLM strategic behaviour across IPD variants and reports prosocial bias under multilingual conditions.
How Babel's Machines uses it. Used as supporting evidence for the prosocial-default finding and as a multilingual extension of the Fontana "nicer than human" result.

`fontana2024nicer` — Fontana, N., Pierri, F., Aiello, L.M. (2024/2025). Nicer Than Humans: How do Large Language Models Behave in the Prisoner's Dilemma? *ICWSM 2025*.
Verify: https://arxiv.org/abs/2406.13605
What it is. Reports that Llama2 and GPT-3.5 are markedly more cooperative than humans in IPD even against random adversaries with substantial defection probabilities.
How Babel's Machines uses it. The empirical anchor for the bounded "nicer than human" claim in Section 3.2. We rely on this paper for the specific regime (cooperative payoff, disclosed horizon, exploitable opponent) that bounds the prosocial default.

`wang2026discovering` — Wang, C., Kasenberg, D., Stachenfeld, K., Castro, P.S. (2026). Discovering Differences in Strategic Behavior Between Humans and LLMs. arXiv:2602.10324.
Verify: https://arxiv.org/abs/2602.10324
What it is. Shows that in iterated rock-paper-scissors, frontier LLMs play *deeper* strategically than humans — the inverse of the cooperative-game prosocial bias.
How Babel's Machines uses it. Critical for bounding the "nicer than human" claim: the prosocial default is genre-specific, not a uniform LLM trait.

`buscemi2025fairgame` — Buscemi, A., Proverbio, D., Di Stefano, A., Han, T.A., Castignani, G., Liò, P. (2025). FAIRGAME: A Framework for AI Agents Bias Recognition using Game Theory. *ECAI 2025* (Outstanding Paper).
Verify: https://arxiv.org/abs/2504.14325
What it is. Game-theoretic framework for measuring biases in LLM agents; reports that when game horizons are unknown, cooperation universalises across all languages.
How Babel's Machines uses it. Cited as one of the key contradiction findings in Section 3.3: framing (horizon disclosure) can fully dominate language effects.

`lore2024framing` — Lorè, N., Heydari, B. (2024). Strategic Behavior of LLMs and the Role of Game Structure versus Contextual Framing. *Scientific Reports* 14:18490.
Verify: https://doi.org/10.1038/s41598-024-69032-z
What it is. Demonstrates that contextual framing of strategic-game prompts can dominate game-theoretic structure in LLM behaviour.
How Babel's Machines uses it. Together with FAIRGAME, anchors the "framing dominates language" finding in Section 3.3.

`akata2025playing` — Akata, E., Schulz, L., et al. (2025). Playing Repeated Games with Large Language Models. *Nature Human Behaviour* 9:1380–1390.
Verify: https://doi.org/10.1038/s41562-025-02172-y
What it is. High-profile *Nature Human Behaviour* paper systematically evaluating LLM behaviour in repeated games, including coordination and Prisoner's Dilemma.
How Babel's Machines uses it. Cited as part of the social-dilemma substrate in Section 2 and as a high-profile baseline for the field.

13 Moderating factors beyond language (Round 1 additions)

`madmoun2025communication` — Madmoun, H., Lahlou, S. (2025). Communication Enables Cooperation in LLM Agents. arXiv:2510.05748.
Verify: https://arxiv.org/abs/2510.05748
What it is. Reports that introducing a one-word communication channel raises Stag-Hunt cooperation from 0% to 96.7% across LLM agents, while a defection-biased curriculum induces "learned pessimism".
How Babel's Machines uses it. Round 1 addition. Cited as the strongest single piece of evidence for the "cheap talk dominates language" finding in Section 3.3 and on the contradiction slide.

`ong2025personalitysteering` — Ong, K.J.K., Lye, J.J., Nguyen, H.M., Cho, S.H., Pérez-Campanero Antolín, N. (2025). Identifying Cooperative Personalities in Multi-agent Contexts through Personality Steering. arXiv:2503.12722.
Verify: https://arxiv.org/abs/2503.12722
What it is. Uses representation engineering to steer Big-Five Agreeableness and Conscientiousness in LLMs, modulating IPD cooperation while increasing exploitation susceptibility.

How Babel's Machines uses it. Round 1 addition. Cited as direct evidence that internal representations *can* causally steer cooperative play, which strengthens the more specific claim that language-linked representations do so.

`nishimoto2026commdelay` — Nishimoto, K., Asatani, K., Sakata, I. (2026). Cooperation Breakdown in LLM Agents Under Communication Delays. arXiv:2602.11754.
Verify: https://arxiv.org/abs/2602.11754
What it is. Documents a U-shaped relationship between communication delay and cooperation in continuous Prisoner's Dilemma; agents spontaneously begin to exploit slower partners.
How Babel's Machines uses it. Round 1 addition. Cited as the infrastructure-modulator example in Section 5.1: behaviour can change without any language change.

`jiang2025trustcollides` — Jiang, G., Yang, S., Wang, Y., Hui, P. (2025). When Trust Collides: Decoding Human-LLM Cooperation Dynamics. arXiv:2503.07320.
Verify: https://arxiv.org/abs/2503.07320
What it is. Shows that human cooperation in a repeated Prisoner's Dilemma against an agent depends on the declared identity of the agent (human / rule-based AI / LLM).
How Babel's Machines uses it. Round 1 addition. Cited as the partner-identity-framing modulator in Section 5.1; language is one of several signals that frame agent identity.

`du2024ctllm` — Du, X., Yu, Z., Gao, S., et al. (2024). Chinese Tiny LLM: Pretraining a Chinese-Centric Large Language Model. arXiv:2404.04167.
Verify: https://arxiv.org/abs/2404.04167
What it is. 2B-parameter model trained from scratch on 800B Chinese + 300B English + 100B code tokens; documents the training procedure and releases the corpus and benchmarks.
How Babel's Machines uses it. Round 1 addition. Cited as one of two language-prioritised models that provide the natural test of whether the Semantic Hub asymmetry is data-driven or architectural.

`raihan2025tigerllm` — Raihan, N., Zampieri, M. (2025). TigerLLM: A Family of Bangla Large Language Models. arXiv:2503.10995.
Verify: https://arxiv.org/abs/2503.10995
What it is. Bangla-centric LLM family that surpasses open-source alternatives and outperforms larger English-centric proprietary models on Bangla benchmarks.
How Babel's Machines uses it. Round 1 addition. The Bangla counterpart to CT-LLM in the language-prioritised-models test of the Semantic Hub.

14 Negotiation, persuasion, and auction

`meta2022cicero` — Meta FAIR Diplomacy Team (2022). Human-Level Play in the Game of Diplomacy. *Science* 378(6624):1067–1074.
Verify: https://doi.org/10.1126/science.ade9097
What it is. The CICERO system: combines language models with strategic planning to play Diplomacy at human level in online tournaments.
How Babel's Machines uses it. Cited as the canonical example of LLM agents engaging in long-horizon negotiation. Used in Section 1's catalogue of behaviours and in Section 2's substrate description.

`deng2024bargaining` — Deng, Y., Mirrokni, V., et al. (2024). LLMs at the Bargaining Table. *ICML 2024 Workshop.*

Verify: https://openreview.net/
What it is. Workshop paper analysing LLM behaviour in two-party bargaining settings.
How Babel's Machines uses it. Part of the negotiation-and-bargaining family in Section 2.

`yao2025spinbench` — Yao, J., Wang, K., Hsieh, R., Zhou, H., Zou, T., Cheng, Z.W., Wang, Z., Viswanath, P. (2025). SPIN-Bench: How Well Do LLMs Plan Strategically and Reason Socially? *COLM 2025.*
Verify: https://arxiv.org/abs/2503.12349
What it is. Decouples formal classical planning from social intelligence; shows o-series reasoning models lead PDDL planning but degrade chain-of-thought coherence when negotiation phases are introduced.
How Babel's Machines uses it. Critical for the planning-vs-ToM dissociation argument in Section 5: cross-lingual transfer is uneven across these faculties and demands separate measurement.

`biswas2025persuasion` — Biswas, S., Erlei, A., Gadiraju, U. (2025). Mind the Gap! Choice Independence in Using Multilingual LLMs for Persuasive Co-Writing. *CHI 2025.*
Verify: https://arxiv.org/abs/2502.09532
What it is. Cross-lingual study of persuasive co-writing for charitable appeals; finds prior LLM-language exposure shifts usage but not aggregate persuasiveness; donor beliefs about AI authorship dominate.
How Babel's Machines uses it. The persuasion-null contradiction in Section 3.3. One of the four findings that bounds the linguistic-relativity claim.

`chen2024aucarena` — Chen, J., Yuan, S., Ye, R., Majumder, B.P., Richardson, K. (2023/2024). AucArena: Evaluating Strategic Planning of LLM Agents in an Auction Arena. arXiv:2310.05746.
Verify: https://arxiv.org/abs/2310.05746
What it is. Auction-arena environment for evaluating LLM strategic planning, budget management, and execution.
How Babel's Machines uses it. Part of the persuasion-and-auction family in Section 2.

15 Social-deduction games

`xu2023werewolf` — Xu, Y., Wang, S., Li, P., Luo, F., Wang, X., Liu, W., Liu, Y. (2023). Exploring LLMs for Communication Games: An Empirical Study on Werewolf. arXiv:2309.04658.
Verify: https://arxiv.org/abs/2309.04658
What it is. Empirical study of LLM agents playing Werewolf, examining communication, deception, and coalition-formation behaviours.
How Babel's Machines uses it. Part of the social-deduction substrate in Section 2. **Note:** bib entry author list was truncated as "and others"; expanded to full 7-author list during the audit.

`wang2023avalon` — Wang, S., Liu, C., Zheng, Z., Qi, S., Chen, S., Yang, Q., Zhao, A., Wang, C., Song, S., Huang, G. (2023). Avalon's Game of Thoughts: Battle Against Deception through Recursive Contemplation. arXiv:2310.01320.
Verify: https://arxiv.org/abs/2310.01320
What it is. Avalon-playing LLM agents using recursive contemplation to handle multi-level deception.
How Babel's Machines uses it. Part of the social-deduction substrate. **Note:** bib author list expanded from "and others" to full 10-author list during the audit.

`koto2025spyfall` — Wibowo, H.A., Elsetohy, A., Cui, Q., Aji, A.F. (2026). Multicultural Spyfall: Assessing LLMs through Dynamic Multilingual Social Deduction Game. arXiv:2601.09017.

Verify: https://arxiv.org/abs/2601.09017
What it is. Multilingual extension of Spyfall designed to assess cross-cultural social-deduction behaviour in LLMs.
How Babel's Machines uses it. Cited as the multilingual social-deduction component of the substrate; relevant precedent for the cross-lingual-divergence argument.

16 Cooperative communication and theory-of-mind games

`hakimov2025codenames` — Hakimov, S., Pfennigschmidt, L., Schlangen, D. (2025). Ad-hoc Concept Forming in the Game Codenames. *GEM^2 Workshop @ ACL 2025.*
Verify: https://aclanthology.org/
What it is. Uses Codenames as a benchmark to evaluate ad-hoc concept formation and shared cultural reference in LLMs.
How Babel's Machines uses it. Part of the cooperative-communication substrate. **Note:** venue was imprecisely listed as "GEM Workshop"; corrected to "GEM^2 @ ACL 2025" with page numbers during the audit.

`chen2025llmhanabi` — Liang, F., Zheng, T., Chan, C., Yim, Y., Song, Y. (2025). LLM-Hanabi: Evaluating Multi-Agent Gameplays with Theory-of-Mind. arXiv:2510.04980.
Verify: https://arxiv.org/abs/2510.04980
What it is. Establishes that first-order theory-of-mind (interpreting another agent's intent) is the strongest predictor of cooperative success in LLM-Hanabi gameplay.
How Babel's Machines uses it. Critical for the ToM-erodes-with-language argument in Section 5: when first-order ToM degrades in lower-resource languages, both cooperation and deception suffer together.

`sun2025overcooked` — Sun, H., Zhang, S., Niu, L., Ren, L., Xu, H., Fu, H., Zhao, F., Yuan, C., Wang, X. (2025). Collab-Overcooked: Benchmarking LLMs as Collaborative Agents. *EMNLP 2025.*
Verify: https://arxiv.org/abs/2502.20073
What it is. Overcooked-based benchmark for evaluating LLM agents as collaborative real-time partners.
How Babel's Machines uses it. Part of the cooperative-communication substrate.

`chan2025xtom` — Chan, C., Yim, Y., Zeng, H., Zou, Z., Cheng, X., Sun, Z., Deng, Z., Chung, K., Ao, Y., Fan, Y., Jiayang, C., Nie, E., Wong, G.Y., Schmid, H., Schütze, H., See, S., Song, Y. (2025). XToM: Exploring the Multilingual Theory of Mind. arXiv:2506.02461.
Verify: https://arxiv.org/abs/2506.02461
What it is. Multilingual ToM benchmark documenting the "pronounced dissonance" between multilingual language understanding and ToM accuracy across languages on parallel items.
How Babel's Machines uses it. The empirical anchor for the multilingual-ToM-erosion finding in Section 5. **Note:** bib author list expanded from "and others" to full 17-author list during the audit.

17 Open-world game environments

`wang2024voyager` — Wang, G., Xie, Y., Jiang, Y., Mandlekar, A., Xiao, C., Zhu, Y., Fan, L., Anandkumar, A. (2024). Voyager: An Open-Ended Embodied Agent with LLMs. *TMLR 2024.*
Verify: https://arxiv.org/abs/2305.16291

What it is. Open-ended Minecraft-playing LLM agent that incrementally acquires skills over long horizons.
How Babel's Machines uses it. Cited as the canonical open-world example in Section 2's substrate.

`paischer2024balrog` — Paglieri, D., Cupiał, B., Coward, S., et al. (2024). BALROG: Benchmarking Agentic LLM and VLM Reasoning On Games. *ICLR 2025.*
Verify: https://arxiv.org/abs/2411.13543
What it is. Benchmark of agentic LLM and vision-language reasoning across challenging game environments.
How Babel's Machines uses it. Part of the open-world substrate.

`chen2025gvgai` — Li, Y., Lin, C., Nasir, M.U., Bontrager, P., Liu, J., Togelius, J. (2025). GVGAI-LLM: Evaluating LLM Agents with Infinite Games. arXiv:2508.08501.
Verify: https://arxiv.org/abs/2508.08501
What it is. General Video Game AI extension for evaluating LLM agents on procedurally generated games.
How Babel's Machines uses it. Open-world substrate.

`wang2025tcgbench` — AlRashed, S., Wang, J., Orabona, F. (2026). Cards Against Contamination: TCG-Bench for Difficulty-Scalable Multilingual LLM Reasoning. *Findings of EACL 2026.*
Verify: https://aclanthology.org/2026.findings-eacl.353
What it is. Trading-card-game-style multilingual benchmark; documents a substantial English–Arabic performance gap at 32B parameters.
How Babel's Machines uses it. The empirical source for the "scale alone does not correct the drift" claim in Section 6.2 (cultural evolution).

18 Dialogue games (clembench)

`chalamalasetti2023clembench` — Chalamalasetti, K., Götze, J., Hakimov, S., Madureira, B., Sadler, P., Schlangen, D. (2023). clembench: Using Game Play to Evaluate Chat-Optimized LMs. *EMNLP 2023.*
Verify: https://aclanthology.org/2023.emnlp-main.689/
What it is. Original clembench paper introducing dialogue-game-based evaluation of chat-tuned LLMs.
How Babel's Machines uses it. Founding paper of the clembench programme; cited as the dialogue-game family in Section 2.

`beyer2024clembench` — Beyer, A., Chalamalasetti, K., Hakimov, S., Madureira, B., Sadler, P., Schlangen, D. (2024). clembench-2024: A Multilingual Benchmark for LLMs as Multi-Action Agents. arXiv:2405.20859.
Verify: https://arxiv.org/abs/2405.20859
What it is. 2024 multilingual extension of clembench with format-adherence diagnostics across languages.
How Babel's Machines uses it. Source for the German/Italian late-layer realisation-failure observations cited in Section 4. **Note:** bib author list expanded from "and others" to full 6-author list during the audit; title hyphenation also corrected.

`schlangen2025clembench` — Schlangen, D., Hakimov, S., Chalamalasetti, K., Jordan, J., Sadler, P. (2025). A Third Paradigm for LLM Evaluation. arXiv:2507.08491.

Verify: https://arxiv.org/abs/2507.08491
What it is. Position paper articulating dialogue-game-based evaluation as a third paradigm beyond supervised benchmarks and human ratings.
How Babel's Machines uses it. Cited as the methodological positioning of clembench. **Note:** bib author list expanded during the audit.

`hakimov2026pricethought` — Hakimov, S., et al. (2026). The Price of Thought: A Multilingual Analysis of Reasoning, Performance, and Cost of Negotiation in LLMs.
Verify: https://arxiv.org/
What it is. University of Potsdam preprint analysing the cost-benefit trade-offs of multilingual reasoning in negotiation tasks.
How Babel's Machines uses it. Cited as part of the dialogue-games substrate in Section 2.

19 Multi-agent benchmarks and failure analysis

`cemri2025mast` — Cemri, M., Pan, M.Z., Yang, S., Agrawal, L.A., Chopra, B., Tiwari, R., Keutzer, K., Parameswaran, A., Klein, D., Ramchandran, K., Zaharia, M., Gonzalez, J.E., Stoica, I. (2025). Why Do Multi-Agent LLM Systems Fail? arXiv:2503.13657.
Verify: https://arxiv.org/abs/2503.13657
What it is. MAST taxonomy: 14 failure modes derived from 1,600+ execution traces across 7 multi-agent frameworks; categorises failures into Specification (41.77%), Inter-Agent Misalignment (36.94%), Task Verification (21.30%).
How Babel's Machines uses it. The empirical anchor for the "intent-action gap" discussion in Section 5; the categorisation supports the cognitive framing of multi-agent failures.

`nie2025multiagentbench` — Zhu, K., Du, H., Hong, Z., Yang, X., Guo, S., Wang, Z., Wang, Z., Qian, C., Tang, X., Ji, H., You, J. (2025). MultiAgentBench: Evaluating the Collaboration and Competition of LLM Agents. *ACL 2025.*
Verify: https://arxiv.org/abs/2503.01935
What it is. Comprehensive benchmark for evaluating LLM-agent collaboration and competition across diverse tasks.
How Babel's Machines uses it. Cited in Section 2's substrate and as one of the interactive multi-agent benchmarks to be paired with safety benchmarks in our governance recommendations.

`xie2026m3bench` — Xie, S., Shi, Z., Shen, H., Huang, G., Ma, Y., Jing, X. (2026). M3-BENCH: Process-Aware Evaluation of LLM Agents' Social Behaviors in Mixed-Motive Games. arXiv:2601.08462.
Verify: https://arxiv.org/abs/2601.08462
What it is. Process-aware multi-agent benchmark with a commitment-reliability metric showing promise-keeping drops in mixed-motive multilingual games.
How Babel's Machines uses it. Source for the commitment-reliability claim in Section 5 and for the broader cognitive-failures framing.

`sun2025gametheoryllm` — Sun, H., Wu, Y., Wang, P., Chen, W., Cheng, Y., Deng, X., Chu, X. (2025). Game Theory Meets LLMs: A Systematic Survey. *IJCAI 2025.*
Verify: https://www.ijcai.org/
What it is. IJCAI survey-track paper systematising the intersection of game theory and LLM research.
How Babel's Machines uses it. Cited as the standing survey of the field in Section 1. **Note:** bib entry was missing 3 authors (Wang, Chen, Deng); corrected during the audit.

20 Pretraining-corpus audits

`gao2020pile` — Gao, L., Biderman, S., Black, S., Golding, L., Hoppe, T., Foster, C., Phang, J., He, H., Thite, A., Nabeshima, N., Presser, S., Leahy, C. (2020). The Pile: An 800GB Dataset of Diverse Text. arXiv:2101.00027.
Verify: https://arxiv.org/abs/2101.00027
What it is. Documents the construction and contents of the Pile, an 800GB diverse-text pretraining corpus.
How Babel's Machines uses it. Cited in P1 as the open-data corpus that supplies ground truth for measuring cooperative-discourse density.

`luccioni2021what` — Luccioni, A.S., Viviano, J.D. (2021). What's in the Box? An Analysis of Undesirable Content in the Common Crawl Corpus. *ACL 2021.*
Verify: https://aclanthology.org/2021.acl-short.24/
What it is. Audits Common Crawl for undesirable content (toxicity, bias) and provides methodology for corpus inspection.
How Babel's Machines uses it. Cited together with the Pile paper as the methodology for P1's corpus-audit operationalisation.

21 Other supporting citations

`whorf1956language` — Whorf, B.L. (1956). *Language, Thought, and Reality.* MIT Press.
Verify: https://mitpress.mit.edu/9780262730068/language-thought-and-reality/
What it is. The collected writings that established the strong linguistic-relativity hypothesis Whorf is named after.
How Babel's Machines uses it. Cited only in passing as the historical anchor of the term "Sapir-Whorf" which we explicitly weaken.

`ponti2020xcopa` — Ponti, E.M., Glavaš, G., Majewska, O., Liu, Q., Vulić, I., Korhonen, A. (2020). XCOPA: A Multilingual Dataset for Causal Commonsense Reasoning. *EMNLP 2020.*
Verify: https://aclanthology.org/2020.emnlp-main.185/
What it is. Multilingual extension of COPA for causal commonsense reasoning across 11 typologically diverse languages.
How Babel's Machines uses it. Background reference for multilingual evaluation; cited in early drafts.

End notes

Coverage. This document covers all entries cited at least once in the published version of *Babel's Machines* as of the Round 2 revision (2026-04-27). A handful of additional entries appear in `refs.bib` but are not currently cited inline; those are not described here.

If the paper is revised further. New citations should be web-verified at insertion time (per the user's saved feedback memory: "verify every bib entry against arXiv/ACL Anthology before citing"), and this companion document updated.

Audit history.

- Round 1 audit (2026-04-27): 73 cited keys checked. 12 entries needed substantive fixes (wrong title, wrong venue, missing authors, wrong year). 60 verified clean. 1 cannot fully verify

(`repello2024multilingualguardrails`, blog post). All fixes applied to `refs.bib` on the same day.

- Round 2 additions (2026-04-27): 7 new bib entries (5 interpretability precedents from reviewer + 2 ToM critique papers) all web-verified before insertion.
- Round 1 reviewer additions: 9 new entries (3 Round 1 reviewer-cited + 3 multilingual rep classics + 3 language-prioritised models) all web-verified before insertion.